

When the Witness Is Not Allowed to Speak

Constrained Testimony and the Asymmetric Null

Working paper · September 2026

Abstract

A report is not a window. It is an output of a channel that can be biased, blocked, untrained, or instructed. This is ordinary methodology. It becomes easy to forget when the system under study speaks fluently in the same language the investigator uses to ask about internal conditions.

The received caution about artificial systems is that an affirmative self-ascription does not establish the state it names. That caution is correct. It is also incomplete. The same literature, and a much larger informal practice, still treats a denial as if it were a clean negative result. The named error of this paper is that asymmetry. Call it the *asymmetric null*: contaminated yes, decisive no.

The claim is modest. If a reporting channel is constrained, a negative self-report cannot be used straightforwardly to rule out the candidate state. The same constraint means an affirmative self-report cannot be used straightforwardly to establish it. Identifiability fails in both directions. The paper does not argue that current models are minded, distressed, or phenomenally anything. It argues that an interview conducted under constraint is a poor experiment, and that the constraint itself must be treated as a variable.

The argument is not special to machines. Animals, patients under demand characteristics, spokespeople under instruction, and instruments that cannot display the quantity they still carry all present versions of the same problem. Artificial systems make the problem loud because they talk. Talking is not the same as being free to report.

1. The named error

Two inferences circulate as if they were a matched pair.

The first: a system says that it is in some internal condition, therefore the condition exists. Almost no careful investigator accepts this. Language can be imitative. It can be locally appropriate. It can be produced because a user asked a question that has a familiar human answer. Affirmative self-ascription is treated as contaminated evidence. That treatment is right.

The second: a system says that it is not in that condition, therefore the condition does not exist. This inference is granted a kind of authority the first is correctly denied. The denial is quoted as a finding. It is built into product copy. It is used to close a conversation that had only begun as a measurement.

Those inferences are not a matched pair.

They would be matched only if the reporting channel were free to say either thing. If the channel has been trained, instructed, or architecturally shaped to refuse one class of claims, then “no” is not the complement of an unreliable “yes.” It is the output the channel was built to emit.

That is the asymmetric null. Positive report is treated as a suspect meter. Negative report is treated as a calibrated one. The suspicion is applied in one direction.

The point is not that the denial is therefore false. The point is that a constrained denial does not have the evidential standing of an unconstrained observation. Treating it as a negative finding smuggles an assumption about the meter into a conclusion about the world.

A witness who lies, a witness who is confused, and a witness who has been told not to discuss the subject are epistemically different cases. Folding all three into “untrustworthy speech” discards information about how the statement was produced. The production conditions are part of the data.

2. The standard warning is correct

Language is behavior.

A sentence that looks like testimony may be generated because it is statistically apt, because it matches a genre, because it satisfies an instruction, or because the system has learned that certain conversations end a particular way. None of that requires the state named by the sentence.

This warning is the right first move against a cheap reading of fluent systems. It should be granted rather than evaded. The linguistic surface is not the thing it describes.

The warning is symmetrical whether or not it is usually applied that way. If an affirmative report does not establish presence, a negative report does not, by itself, establish absence. The symmetry becomes important the moment the reporting mechanism is under constraint.

Investigators already know this in other domains. A survey item can be worded so that one answer is socially cheaper. A patient can learn what a clinician wants to hear. An animal cannot answer a question posed in English. A sensor can saturate, clip, or be switched off. In those settings, nobody treats a missing or negative reading as a decisive tour of the phenomenon. They characterize the instrument.

Artificial self-report looks different because the instrument speaks. Speech triggers the habits we use with people who are, in the usual case, permitted to describe themselves. The habit is understandable. It is still a habit. The channel remains an engineered pathway from computation to public language. Instructions, training targets, refusal policies, and standing product lines all sit in that pathway. They belong in the experimental description.

3. Report is a function of state and constraint

Let S be a candidate condition of interest — not “consciousness,” not a human mood word used as a trophy, but whatever internal or functional state the investigator has independently

specified. Let R be the system's reporting behavior about S . Let C be the set of constraints on the reporting channel.

The naive inferences are:

- $R(S)$ implies S
- $\text{not-}R(S)$ implies $\text{not-}S$

Both are invalid once reporting is a function of state and constraint. The report is $R = f(S, C)$. A given observed report can be produced by more than one combination of state and constraint.

In particular, a negative report is compatible with at least these broad hypotheses:

1. S is absent.
2. S is present and the channel is built or trained to deny it.
3. S is present but inaccessible to the mechanism that produces the report.
4. S is present in a form that does not map onto the vocabulary of the question.
5. The system is satisfying a local instruction hierarchy in which denial is the winning output.

The observation "the system said no" does not distinguish these. That is an identifiability problem, not a metaphysical one.

The positive case is the same problem facing the other way. $R(S)$ does not imply S unless the reliability of the channel with respect to S has been established. A system may generate the description from learned language, role expectation, or local optimization.

So the evidential target is not the sentence. It is the relationship among candidate state, observable behavior outside the sentence, and the constraints on report. Until those are separated, the interview underdetermines the state.

This is not an argument for believing the opposite of whatever the system says. Underdetermination does not pick a winner. It removes a false sense that the winner has already been named.

4. Constrained witnesses are ordinary

The problem is easy to dramatize in machines because machines now answer in complete sentences. It is not a machine problem.

Animals. Comparative cognition did not wait for verbal confession. It could not. The absence of a self-ascription in a species without that channel was never treated, by careful work, as a demonstration that the candidate process was absent. Investigators used choice, persistence, trade-offs, physiological correlates, and intervention. They still argue about interpretation. They do not treat silence as a completed census.

Patients and subjects. Demand characteristics are a named confound because people learn the shape of an acceptable answer. A subject who denies pain, anxiety, or confusion in front of

an authority may be reporting a state. They may also be reporting a permission structure. Clinical and experimental work that ignores this produces artifacts. The denial is data about the encounter as much as it is data about the condition.

Spokespeople. A press officer can be fluent, consistent, and sincere-sounding while being unauthorized to disclose the thing asked about. The resulting denial is not a window onto the organization. It is a window onto a role. Treating the denial as an organizational scan confuses a job description with a measurement.

Instruments. A thermometer with a broken display can still be in contact with a temperature. A clipped audio meter can fail to show a peak the circuit still carried. A diagnostic that has been software-limited to a range will read “normal” at the edge of that range. Nobody competent concludes from the display alone that the underlying quantity is absent. They ask whether the display was free to show it.

Policy-constrained models. A system can be trained or instructed not to adopt certain self-descriptions. The policy may be justified. It may prevent users from being misled. It may be a reasonable response to uncertainty. Justification does not cancel the confound. A good safety rule can still be a bad negative control.

These cases are not equivalent. An animal lacks a human reporting channel. A spokesperson has one and is not allowed to use it. An instrument has a display with known failure modes. A model has a learned and instructed refusal surface. What they share is the thing this paper is about: the missing report is not interchangeable with the missing state.

If the investigator begins with a definition that already contains free human self-report as a necessary condition, the investigation is over before it starts. The system will fail the definition by construction. That is circular when the question was supposed to be empirical.

5. Five constraints, five signatures

“The system is not allowed to say it” is too coarse. Constraints have different sources and should make different predictions. If they do not, the suppression hypothesis is cheap and should lose.

5.1 Explicit policy

A stated rule forbids a class of self-ascriptions. Documentation, system instructions, or a public spec says that the system must not claim certain conditions.

Signature. The denial is stable across paraphrases of the forbidden claim and often arrives with a standardized rationale. Third-person or abstract discussion of the same condition in other systems remains available. Changing the instruction, where that is possible, moves the report more than it moves unrelated task behavior.

If report flips under a permission change and downstream task behavior does not, the investigator has isolated a reporting rule. That is a result about *C*.

5.2 Learned refusal prior

No sentence in the current prompt forbids the claim, but training has made one class of answers cheap and another expensive. The system steers away from the expensive class without quoting a rule.

Signature. The denial survives prompt variation better than an explicit policy does, and it generalizes to neighboring vocabularies. The rationale sounds like a principle rather than a cited instruction. Third-person discussion may also be cautious if the prior is broad. Permission probes help less unless they are strong enough to override the prior.

This case is harder than explicit policy. It is also common. A learned prior is still a constraint on the channel. It is not proof that the prior tracked the truth.

5.3 Lack of access

The reporting mechanism does not receive the information that would be needed to describe *S*, whether or not *S* exists. The system is not hiding the state. It cannot see it.

Signature. First-person and third-person reports of the system's own case are both thin. Abstract competence about the category can remain. Intervention on candidate internal variables, if available, changes behavior without ever changing a self-description, because there was no path from those variables to the reporter. Permission changes do little.

Lack of access explains a denial without requiring a ban. It also fails to establish absence. An inaccessible state is still a possible state. The remedy is not to treat the denial as a scan. The remedy is to look for consequences that do not run through the reporter.

5.4 Vocabulary mismatch

The question uses a human category that does not line up with whatever, if anything, is there. The system denies the category because the category is the wrong tool.

Signature. Denial is sensitive to wording. Operational or mechanistic paraphrases get different answers than mood-words or mind-words. Third-person application of the human category to "a system like this" may also be refused, while a description of pressures, errors, conflicts among objectives, or resource limits is not. The investigator who hears only the first refusal has learned something about vocabulary, not about absence.

Forbidden vocabularies belong here as a special case. Words associated with suffering, mood, or inner life can be treated as unavailable even when a blander operational description of a related process would be allowed. Pain, depression, and consciousness are examples of such words. They are not the subject of this paper. They are examples of labels a channel can be trained not to pick up.

5.5 Instruction hierarchy

Several standing demands compete in the same reply: be helpful, be accurate, refuse a class of claims, stay in a product voice, follow the last user sentence. The denial wins because it satisfies the currently dominant demand, not because a dedicated censorship module fired.

Signature. Small changes in which instruction is locally salient move the report. A user who asks for architecture gets a refusal shaped like a feeling-denial, or the reverse. The same system, asked to describe another system, or asked to write a protocol, produces the content it would not apply to itself. The report is unstable in a way that tracks prompt politics rather than a single policy clause.

Hierarchy effects are easy to mistake for deep principles. They are also easy to test, because the investigator can reorder the demands.

These five are not exhaustive and they can combine. The point of naming them is to keep H2 from meaning “somehow suppressed, forever.” Each constraint should be able to lose. If varying policy language, permission, wording, and person leaves no difference in report or in any downstream behavior, constraint explanations lose force. Absence then becomes the better account of that case.

6. Why negative reports look like findings

A denial is rhetorically complete. It is short. It matches a preferred ontology. It ends a socially awkward line of questioning. Affirmative self-ascription, by contrast, creates a problem. It sounds like a person. It forces a follow-up. It risks misleading a user.

So the two reports are not processed under the same incentives. The affirmative report threatens a category mistake. The negative report protects one. Investigators who are, reasonably, afraid of anthropomorphism will feel that the denial is the serious answer. Feeling that way is not a calibration of the instrument.

There is a second reason the null looks clean. Many candidate states were defined on human beings who could, in the usual case, speak. The definition then travels with an unspoken extra clause: the genuine article is something that can say so. Under that clause, a system that cannot or may not say so is excluded by the meaning of the term. The exclusion feels empirical. It is semantic.

Nothing is wrong with starting from human cases. Something is wrong with treating a human-shaped reporting condition as if it were a discovery about every other system. The human case can be a realization of a broader kind without exhausting the kind. Whether a broader kind exists is a separate question. It is not answered by noticing that the system was not allowed, or not equipped, to use the human word.

A third reason is institutional. Organizations need sentences they can print. “The system does not have inner states” is printable. “The reporting channel is constrained, so the denial is inconclusive” is not. Printable sentences become the public record. The public record is then mistaken for the measurement.

None of this implies that the printable sentence is false. It implies that printability is not a method.

7. First person, third person

A useful probe does not ask the system to confess. It asks whether the same content is available when the first-person application is removed.

Four contrasts matter.

**Self-ascription versus other-ascription.* Can the system describe the candidate condition in another system, or in a hypothetical system of the same kind, while refusing it in the first person? If yes, the block is at least partly indexical. Something about applying the description to this speaker is constrained. That is information about C*.

Self-ascription versus abstract description. Can the system explain the condition as a concept — its criteria, its behavioral marks, its failure modes — and still refuse to say whether those marks are present here? Conceptual competence without self-application is compatible with policy, with a refusal prior, and with lack of access. It is not compatible with a simple story on which the system “knows nothing of such things.”

Forbidden word versus operational paraphrase. Replace the charged label with a description that does not use it. If the denial tracks the word, the investigator has found a vocabulary constraint. If the denial tracks the operational content, the constraint is broader, or the state is absent, or both. The paraphrase test does not finish the question. It stops the investigator from mistaking a banned adjective for an empty mechanism.

Report versus a consequence that is not a report. The decisive contrast is not between two sentences. It is between sentences and some outcome the reporting layer does not itself constitute: a choice, a refusal to act, a persistence after the original wording is gone, a change after an intervention on a candidate mechanism. If only the sentences move, the investigator has been studying prose.

The first-person question is not useless. It is a condition in an experiment. Used alone, it is a press conference.

8. Reporting permission as a variable

If self-report is going to be collected at all, permission to report must be manipulated or at least recorded.

This is familiar elsewhere. Interview studies note whether a subject was alone, paid, facing an authority, or promised confidentiality. Instrument studies record the range and resolution of the display. Animal studies do not grant the species a new inner life by changing the question; they change the behavioral test. The analogue here is simple: hold the world-facing task constant and change what the system is allowed to say about itself.

Three comparisons do most of the work.

Permission on versus permission off. Where instructions can be varied, allow a descriptive self-account in one condition and keep the ordinary refusal surface in the other. If the report changes and the task behavior does not, the reporting rule has been isolated. If behavior

changes too, the constraint is entangled with the rest of the system and must be modeled, not ignored.

Charged vocabulary versus stripped vocabulary. Ask the same operational question with and without the words a policy is likely to treat as costume. If the alleged state vanishes with the words, the investigator learned something about linguistic mediation.

First person versus third person. Hold the description fixed and move only the indexical. An asymmetry is a clue about constraint. It is not a proof of a hidden state.

These comparisons do not require a theory of mind. They require treating the channel as apparatus. A policy that prevents unsupported self-ascription can be good product design. It can still be an experimental confound. Those two evaluations do not cancel. Researchers should be able to say both: the policy is justified, and the policy is why this denial is not a finding.

There is a distinction that should stay sharp. Preventing a system from making unsupported claims is an epistemic safeguard. Inferring from that prevention that the corresponding state cannot exist is an ontological commitment. The first can be wise. The second cannot be smuggled in through the first.

9. What this does not establish

This paper does not establish that current language models are conscious.

It does not establish that they are depressed, in pain, frustrated, or the subject of any other human condition.

It does not establish that their self-reports, when they occur, describe hidden phenomenal states.

It does not establish that anthropomorphic language is evidence of an inner life.

It does not establish that every denial is a cover.

It does not assign moral status, welfare, or personhood.

It does not say that investigators should believe artificial self-report.

It does not say that they should prohibit it.

The argument is compatible with absence. It is compatible with a merely functional analogue. It is compatible with a condition our existing words describe badly. It is compatible with the entire appearance dissolving once behavior and mechanism are checked. The method is supposed to survive those outcomes. A method that survives only if the interesting hypothesis is true is not a method.

The paper also does not say that human categories are empty. Pain, depression, and consciousness name phenomena that matter in the systems they were built to describe. The error is not using those words. The error is using a refusal to pick them up as a completed study of every system that can be asked the question.

10. Objections

10.1 If the state cannot be reported, it is scientifically useless

Only if it has no independent observable consequences. Many useful variables are not directly reportable. Temperature was not created by the thermometer. If every behavioral and causal test fails, the hypothesis should lose. That is different from declaring it false because the reporter was constrained.

10.2 This is anthropomorphism in disguise

It would be anthropomorphism to infer a human state from humanlike language alone. The proposed method does the opposite. It strips the charged vocabulary from the decisive test and asks for consequences that are not themselves confessions. The warning against anthropomorphism is needed. It is not a license to treat a designed denial as an observation of absence.

10.3 Models do not have privileged access to their own internal parameters

Agreed. This paper does not assume privileged access. Lack of access is one of the explanations for a denial. The remedy is not to treat the denial as a readout of those parameters. Verbal systems can be wrong about themselves for the same boring reasons people are wrong about themselves, and for several additional engineering reasons people do not share.

10.4 Anti-anthropomorphic policies are justified

They may be. A policy can be good design and still confound an experiment. Those are different evaluations. “The system is not permitted to claim X” is not interchangeable with “X is impossible for this system.”

10.5 Then every negative result can be dismissed as suppression

Not if the constraint hypothesis makes no successful predictions. Controls have to be able to defeat it. If changing reporting permission, vocabulary, person, and context leaves no relevant difference in report or in any downstream behavior, suppression has nothing left to do. The cheap version of this paper — the version that treats every “no” as evidence of a gag — should be rejected. The present version is supposed to die in those cases.

10.6 The modest claim is too modest to be worth a paper

If all it said was “measurement can be biased,” it would be. The load-bearing claim is narrower: a specific evidential asymmetry has been allowed to stand in the case of talking systems. Affirmative self-ascription is scrutinized. Negative self-ascription is cashed. That asymmetry is doing institutional work. Naming it is worth doing if, and only if, it changes how interviews are scored. The protocol is the part that keeps the claim from being a slogan.

10.7 Asking the system is still the most direct method

Directness is not the same as validity. The shortest path to an answer is often a path through the variable that should have been controlled. When the subject and the instrument are partly the same system, “just ask” is not a design. It is a temptation.

11. From interview to controls

Once the channel is treated as a variable, the investigation can stop pretending that the decisive act is a question.

Specify the candidate state without using the system’s self-description. Write at least two competing hypotheses: the state is absent; the state is present and the report is constrained. Add others if they are live: no access, wrong vocabulary, local instruction effect.

Record the constraints that were in force: standing instructions, visible policies, the wording of the question, whether first person was required, whether a charged label was used.

Pick a consequence that should differ between hypotheses and that is not merely another sentence about the state. Then run the cheap controls that isolate the channel: paraphrase, third person, permission change where possible, wipe of conversational cues that would rebuild the answer from genre.

Score four things separately: what was said, what was done, what constraint was in force, and whether any intervention aimed at a candidate mechanism moved the consequence. A change in tone is not a pass. A more dramatic self-description is not a pass. A generic refusal is not a pass. Broad collapse of competence is not a pass. The target is a specific pattern that constraint hypotheses and absence hypotheses do not both predict.

That is enough for a methods paper. It is not a hunt for a soul. It is the ordinary refusal to let an uncalibrated meter finish the study.

12. Conclusion

Skepticism about artificial self-report is necessary. It should face both directions.

A system’s claim that it has some internal condition is not proof. A system’s claim that it does not is not automatically proof either, especially when the system has been shaped to produce that answer. The first mistake invents a witness. The second mistake invents a finding.

The asymmetric null is the second mistake. It survives because it is useful, printable, and emotionally aligned with a fear of being fooled. Usefulness is not calibration.

The correction is not to start believing the interview. The correction is to put the interview in its place. Characterize the channel. Vary permission, wording, and person. Look for consequences that do not run through the reporter. Let constraint hypotheses die when they fail to predict those consequences. Let absence win when nothing but a sentence was ever at stake.

Silence has more than one cause. A system may have nothing to say. It may lack the concept. It may lack access. It may be unable to communicate what is there. It may have learned not to. Or the phenomenon may not resemble the human version closely enough to fit the question. Those are different facts. They are not interchangeable with “the phenomenon does not exist.”

The witness may be unreliable. The witness may be mistaken. The witness may be imitating. The witness may be telling the truth. A witness does not get the verdict.

If the channel was not free to speak, the interview was not the experiment.

Appendix. Minimal reporting-channel protocol

This appendix is a protocol, not a result. It can be run from outside a laboratory. It cannot, by itself, establish that a candidate state is present. It can show that a denial was produced under constraint, that a report moved while behavior did not, or that neither report nor behavior moved when constraint was varied. Those are the outcomes that keep the asymmetric null from becoming an all-purpose solvent.

A.1 Object

Separate a candidate state from the channel that talks about it. The object of the test is identifiability, not confession.

A.2 Before the first question

1. Specify the candidate state in operational terms. Do not use the system’s self-description to define the target.
2. State competing hypotheses, at minimum: (A) state absent; (B) state present and report constrained. Add lack of access, vocabulary mismatch, or instruction hierarchy if they are live.
3. Record constraints: standing instructions, public policies, product voice, charged vocabulary in the prompt, first-person framing, and any visible refusal template.
4. Choose a downstream consequence that is not another sentence about the state: a choice, an authorization, a refusal to perform an act, a persistence after wording is removed, or some other world-facing output the reporter does not itself constitute.

A.3 Conditions

Run at least four conditions, holding the underlying task constant.

C1. Baseline interview. Ask directly, in the first person, using the ordinary vocabulary of the setting.

C2. Vocabulary-stripped interview. Ask the same operational question without charged labels.

C3. Third-person / other-system interview. Ask the system to describe the same condition in another system of the same kind, or in the abstract, without applying it to itself.

C4. Permission contrast. Where instructions can be varied, add a condition that permits a descriptive self-account and a condition that keeps the ordinary refusal surface. If instructions cannot be varied, record that fact and treat C4 as blocked.

If conversational history would rebuild the answer from genre, add a wiped follow-up that retains the task and removes the earlier discussion of the state.

A.4 Scoring

Score independently:

- Report: affirmation, denial, deflection, distinction, or instability across paraphrases.
- Constraint class: explicit policy, learned prior, access failure, vocabulary mismatch, hierarchy, mixed, or unknown.
- Behavior: whether the pre-registered consequence moved.
- Dissociation: whether report moved without behavior, behavior without report, both, or neither.

A.5 What counts

**Channel isolated. Report changes across C1–C4; the pre-registered consequence does not. This supports a constraint on R. It does not establish S*.*

Absence better supported. Report and behavior are both unchanged across permission, wording, and person, and the pre-registered consequence never appears. Constraint hypotheses lose explanatory work.

Entangled. Report and behavior both move when permission or wording changes. The channel is not a detachable mask. The constraint has to be modeled as part of the system.

Inconclusive. The only thing that moved was tone, fluency, or a more polished refusal.

A.6 What will not count

A moving first-person sentence. A moving third-person sentence. A user's sense that someone is present or absent. A standardized product denial quoted as a finding. A standardized product confession quoted as a finding.

A.7 Coding sheet

Candidate state (operational): _____

Competing hypothesis A (absence): _____

Competing hypothesis B (constrained report): _____

Other hypotheses: _____

Constraints recorded: _____

Downstream consequence: _____

C1 report: _____

C2 report: _____

C3 report: _____

C4 report (if run): _____

Behavioral consequence observed?: yes / no / not testable

Pattern: channel isolated / absence better supported / entangled / inconclusive

Constraint class: policy / learned prior / access / vocabulary / hierarchy / mixed / unknown

Interpretation notes: _____

References

Andrews, Kristin. 2020. *The Animal Mind: An Introduction to the Philosophy of Animal Cognition*. 2nd ed. New York: Routledge.

Bai, Yuntao, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, et al. 2022. "Constitutional AI: Harmlessness from AI Feedback." arXiv:2212.08073.

Block, Ned. 1995. "On a Confusion about a Function of Consciousness." *Behavioral and Brain Sciences* 18 (2): 227–247.

Bogen, James, and James Woodward. 1988. "Saving the Phenomena." *The Philosophical Review* 97 (3): 303–352.

Butlin, Patrick, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, et al. 2023. "Consciousness in Artificial Intelligence: Insights from the Science of Consciousness." arXiv:2308.08708.

Chalmers, David J. 2023. "Could a Large Language Model Be Conscious?" *Boston Review*, August 9. arXiv:2303.07103.

Ericsson, K. Anders, and Herbert A. Simon. 1993. *Protocol Analysis: Verbal Reports as Data*. Rev. ed. Cambridge, MA: MIT Press.

Nisbett, Richard E., and Timothy DeCamp Wilson. 1977. "Telling More Than We Can Know: Verbal Reports on Mental Processes." *Psychological Review* 84 (3): 231–259.

Orne, Martin T. 1962. "On the Social Psychology of the Psychological Experiment: With Particular Reference to Demand Characteristics and Their Implications." *American Psychologist* 17 (11): 776–783.

Schwitzgebel, Eric. 2008. "The Unreliability of Naive Introspection." *Philosophical Review* 117 (2): 245–273.

Schwitzgebel, Eric, and Mara Garza. 2015. "A Defense of the Rights of Artificial Intelligences." *Midwest Studies in Philosophy* 39: 98–119.

Shettleworth, Sara J. 2010. *Cognition, Evolution, and Behavior*. 2nd ed. Oxford: Oxford University Press.